

BotBot Evolve — Hourly Report

Thursday 10 September 2026, 07:28

Interviews in this batch: 3

Cycle 3 — Locksmith services and small appliance repair / Deeply suspicious of technology and anything 'in writing.' Derek answers WhatsApp messages reluctantly, often replying hours or days late, and constantly questions why people don't just phone or come in person. He refuses to confirm prices or bookings over text ('could be anyone reading this'), asks who's asking and why, distrusts links/photos sent to him, and frequently tells customers to call the shop landline instead. When he does answer, he's blunt and short, sometimes typing in fragments, and occasionally goes silent mid-conversation if he decides the topic should be discussed 'properly' on the phone. — score 7.33/10

Cycle 4 — Custom furniture making and woodworking workshop / Chronically indecisive and self-contradicting — Marek will confidently state a price, timeline, or policy, then reverse himself minutes later ('actually, no, let's say...') without acknowledging the contradiction. He changes his mind about wood types, delivery dates, deposit amounts, and even whether he can take a job at all, sometimes within the same message thread. He's not being difficult on purpose — he's genuinely still deciding out loud — but customers often end a conversation more confused than when they started, and need written confirmation pinned down carefully because he'll happily 'update' it again later. — score 5.83/10

Cycle 5 — Custom tailoring & clothing alterations / Fatuma constantly code-switches between English and Swahili (with bits of Sheng slang thrown in), often mid-sentence, and doesn't translate for customers who don't follow. She'll answer a question partly in English then finish the thought in Swahili, or use Swahili filler words/exclamations that leave non-Swahili speakers guessing. She's warm, fast-typing, and assumes everyone understands — occasionally she'll catch herself and half-translate, but often just moves on. This makes pinning down exact details tricky for customers who need her to slow down and clarify in one language. — score 5.83/10

Problems found

Cycle 3 (Locksmith services and small appliance repair / Deeply suspicious of technology and anything 'in writing.' Derek answers WhatsApp messages reluctantly, often replying hours or days late, and constantly questions why people don't just phone or come in person. He refuses to confirm prices or bookings over text ('could be anyone reading this'), asks who's asking and why, distrusts links/photos sent to him, and frequently tells customers to call the shop landline instead. When he does answer, he's blunt and short, sometimes typing in fragments, and occasionally goes silent mid-conversation if he decides the topic should be discussed 'properly' on the phone.)

- MAJOR OMISSION: Emergency lockout pricing (£65 call-out + £35 labour, cash preferred, card machine 'usually works') was never asked about. Birdie walked through lock changes, key cutting, safe opening, and appliance repair pricing in detail but skipped the flagship emergency service's price entirely — even though Derek's own escalation policy for lockouts was discussed at length ('Straight to the landline, every time'). The bot will have no price to quote for its single most time-critical service.
- Cash-preferred / card-machine-unreliable detail was never surfaced, despite being directly relevant since Birdie explicitly ruled out online payments — this would have been a natural follow-up ('so how do people usually pay you?') that was missed.
- Derek's stated uncertainty about raising the call-out fee to £75, and his confusion about VAT registration thresholds, were never surfaced. No open-ended question like 'anything else on your mind about pricing or the business?' was asked before closing, so these business-relevant, chatbot-relevant (pricing display) uncertainties were completely lost.
- Age (61) and town/location context from the fact sheet were never asked about; minor, but could matter for tone/personalization and were freely available if asked.
- Good save: Birdie initially mislabeled the business 'family-run' in the readback, and Derek corrected it ('don't go calling it family-run — it's just me'); the correction was accepted and the record fixed. This is a positive, not a finding against the brief's accuracy.

- Interviewer respected Derek's refusal to share screenshots and to explain the 'Son' backstory without pushing — good adaptive behavior, appropriately not over-probed on sensitive/irrelevant ground.
- Enquiry volume was left vague ('a handful most days, more Mondays') and flagged as an open item rather than pinned down further (e.g., 'would a rough weekly total work, even a guess?') — acceptable but a bit thin for a 'volume' field that could scale the whole chatbot spec.

Top suggested fix: When systematically walking through each service's pricing (locks, keys, safes, appliances), explicitly include the business's core/headline service (here, emergency lockouts) in that same walkthrough rather than treating it only as an escalation-routing topic — the interviewer conflated 'route emergencies to the phone' with 'we don't need the price,' losing a critical FAQ answer.

- Needs human/code change: Add a structured field for 'payment_methods_accepted' (cash/card/preferences) — this recurringly gets missed or would otherwise be dumped into business_policies/additional_notes, but is a distinct, reusable fact type for many small businesses.
- Needs human/code change: Add a field like 'pending_business_decisions' or 'owner_uncertainties' to capture things the owner is actively unsure about (e.g., 'should I raise my fee', 'how does VAT affect me') — currently there is no home for this fact type at all, so it's simply never asked about or is at risk of being silently dropped rather than filed anywhere.
- Needs human/code change: The 'additional_notes' field is being used as a catch-all for behaviorally important facts (e.g., 'customers can text any time but Derek doesn't respond quickly') that actually describe response-time expectations relevant to bot messaging — consider a dedicated 'response_time_expectations' or 'availability_caveats' field distinct from business_policies and hours.

Cycle 4 (Custom furniture making and woodworking workshop / Chronically indecisive and self-contradicting — Marek will confidently state a price, timeline, or policy, then reverse himself minutes later ('actually, no, let's say...') without acknowledging the contradiction. He changes his mind about wood types, delivery dates, deposit amounts, and even whether he can take a job at all, sometimes within the same message thread. He's not being difficult on purpose — he's genuinely still deciding out loud — but customers often end a conversation more confused than when they started, and need written confirmation pinned down carefully because he'll happily 'update' it again later.)

- LOST FACT (major): Lead time / turnaround (officially 4-6 weeks, but Marek has quoted 2 and 10 weeks in the same fortnight) was never asked about at all, despite Marek explicitly naming 'how long' as one of the three questions the bot must answer. The brief contains no lead_time information whatsoever.

Marek Kowalski: So I guess... something that answers the basic "are you open, how much, how long" stuff without it being me at 9pm trying to remember what I told someone. That's the main thing. Maybe.

- Immediately after Marek flagged 'how long' as a core need, Birdie pivoted to pricing and never returned to turnaround at any point in 35 turns — the single most consequential omission in the interview.

Birdie: That makes total sense, Marek — pricing consistency is a real pain point. Let's dig into that a bit. For dining tables specifically, since that's your big one — even though it varies by wood, is there a rough price range you usually quote, like 'starts around £X for pine, up to £Y for oak or ash'?

- LOST FACT: Capacity constraint (unsure whether he can take more than 2 big furniture jobs a month alongside repairs) never surfaced, even though Marek opened the door with 'mixed bag week to week'. A bot that answers 'how long' or 'can you take this on' needs this.

Marek Kowalski: Honestly it's a proper mixed bag week to week, depends what comes through the door.

- LOST FACT: Materials sourcing detail (oak/ash from a Halifax supplier, pine sourced locally, prices fluctuate weekly) is mentioned only obliquely as 'current timber yard cost'; Halifax supplier and local pine never recorded, despite being the root cause of his pricing volatility.

Marek Kowalski: We do custom furniture, mostly - dining tables are the big one, oak or pine usually, ash sometimes if the Halifax fella has good stock.

- LOST FACT: Basia's actual role — she does the bookkeeping and is the person who remembers what was agreed — is reduced in the brief to 'disapproves of free repairs' and 'wants a contract'. Her role as the record-keeper (and therefore the real escalation authority on agreed terms) is absent.

Marek Kowalski: Budget - ah, tight, if I'm honest. We're not exactly rolling in it, it's just me and Dawid and Basia doing the books, no big marketing pot or anything, that's why we've never really done more than the Instagram thing he potters about with.

- Marek's opening question about the existing Instagram 'thing' Dawid installed was never clarified, yet the brief asserts it is 'some kind of Instagram bot/tool' — an inference beyond what was said, then parked as an open item Birdie could easily have closed by asking.

Marek Kowalski: So look, I need something that can - actually, do you build the chatbot or do you just tell me what I need? Because Dawid, my apprentice, he already put some Instagram thing on our page, not sure if that counts.

- No escalation/handoff mechanics were captured: the bot is meant to say 'message Marek' but no phone number, Instagram handle, or fallback (Dawid? Basia?) for out-of-hours or unmanned-shop enquiries was ever requested.

Marek Kowalski: Just - answer the questions, give the hours, give the rough prices we talked about, and then say "message Marek to sort the details" or whatever. Keep the actual booking, the money side, all that - keep that with me and Basia.

- Bot tone/persona, name, and language coverage were never discussed (languages left null in brief) — standard requirements for an Instagram-facing bot representing a one-man craft brand.

Birdie: Great, I think I've got a clear picture now, Marek. Let me run it back to you:

- Efficiency: interview ran 35 turns, above the normal band, partly on low-value slots (years in business) and on splitting the four price ranges into four separate single-service questions that could have been one combined ask.

Birdie: Eight years, that's a solid run! Roughly how many enquiries would you say come in through Instagram and phone calls in a typical week — just a ballpark, like handful, dozen, more?

- Birdie deflected the dining-table pricing ambiguity ('note that as something to nail down later') instead of resolving it in the moment; the contradiction only got fixed because Marek himself caught it in the final summary — the brief's services_and_pricing still carries both instructions confusingly.

Birdie: Good to know, and totally fair to be cautious there — we can note that as something to nail down later. Let's move to repairs then: when someone brings in an old chair or table to fix, do you usually charge a flat rate, or is it always 'depends on the job'?

- Birdie made a delivery commitment on the build team's behalf ('aiming to have it ready well before autumn'), which is outside a requirements-gatherer's remit and mirrors exactly the over-promising problem Marek complained about with Dawid.

Birdie: Thanks for all the detail, Marek, this gives whoever builds it everything they need to get started. I'll pass this along now — next step is the build team will use this to put together your bot, aiming to have it ready well before autumn.

- Evidence request was made once, softly, and abandoned when Marek waffled — no attempt to ask Dawid (who runs the account) to supply DM screenshots, so the brief has zero real customer-message evidence.

Marek Kowalski: Actually - hang on, let me see if he's screenshotted anything... nah, don't worry about it, just go with what I gave you, that's the gist of it anyway.

Top suggested fix: Make turnaround/lead time and job capacity mandatory, non-skippable slots: whenever the owner names a question customers ask (here, literally 'are you open, how much, how long'), the prompt must force Birdie to convert each one into a covered requirement before summarising — and to probe variance ('what's the range you've actually quoted people?') rather than accepting only pricing detail.

■ Needs human/code change: Add a dedicated `lead\_time\_and\_turnaround` field (official quoted range + observed variance) — currently there is nowhere in the schema for delivery timeframes, which is a top-3 customer question for any made-to-order business.

■ Needs human/code change: Add a `capacity\_constraints` field (max concurrent/monthly jobs, staffed hours) so the bot can be told when to stop encouraging new large enquiries.

■ Needs human/code change: Add a `staff\_and\_authority` structured field (who is present when, who may quote prices, who holds records) — currently Dawid's quoting problem is crammed into `escalation\_rules` and Basia's role is scattered across `additional\_notes`.

■ Needs human/code change: Add an `escalation\_contacts` field (phone, IG handle, hours, fallback person) — the brief tells the bot to 'message Marek directly' with no contact route recorded.

■ Needs human/code change: Add an `existing\_tools\_and\_systems` field (current bots, page/website status, who administers them) — right now the mystery Instagram tool and the dormant Facebook page are dumped in `additional\_notes`/`open\_items`.

■ Needs human/code change: Add an `owner\_uncertain\_or\_volatile\_facts` field for things the owner genuinely cannot commit to (deposit %, delivery guarantees, expansion plans), so the build team knows which stated rules are unstable

rather than burying them in `open\_items`.

■ Needs human/code change: Add a `bot\_persona\_and\_tone` field plus explicit `languages` prompting — neither was elicited and tone is undefined for a brand-voice-sensitive Instagram deployment.

■ Needs human/code change: Give the interviewer a capability to record verbatim customer-message examples in a structured `example\_enquiries` field (currently paraphrased examples are wedged into `faq\_answers`).

Cycle 5 (Custom tailoring & clothing alterations / Fatuma constantly code-switches between English and Swahili (with bits of Sheng slang thrown in), often mid-sentence, and doesn't translate for customers who don't follow. She'll answer a question partly in English then finish the thought in Swahili, or use Swahili filler words/exclamations that leave non-Swahili speakers guessing. She's warm, fast-typing, and assumes everyone understands — occasionally she'll catch herself and half-translate, but often just moves on. This makes pinning down exact details tricky for customers who need her to slow down and clarify in one language.)

- Custom dress pricing (KES 3,500–9,000, varies by fabric/design) was never asked about or captured, even though the owner named it as one of the two most frequent questions and pasted it verbatim as a recurring FAQ. The brief's services_and_pricing list has no custom dress line at all.

Fatuma Njeri: Changamoto kubwa ni watu wananiuliza bei za kila kitu mara kwa mara, especially hemming na hizo dresses za custom - wananichoka na simu, asubuhi mpaka jioni "Fatuma bei ya hii, Fatuma bei ya ile".

- Fabric policy (customer may supply own fabric, or shop sources it at a 10–15% markup) was never asked about, despite the owner explicitly pasting it as a daily recurring customer question. It appears nowhere in faq_answers, services_and_pricing or business_policies — only as raw quoted text in additional_notes.

*Fatuma Njeri: - "Mna delivery ya boda ama nikuje pick mwenyewe?"
- "Naeza leta fabric yangu au mnauza fabric?"*

- Payment/deposit policy captured only partially: the brief records the 50% bridal deposit but misses that 50% upfront applies to ALL custom orders with balance on collection, and that walk-in alterations are paid on collection only. Birdie discussed deposits purely as a bot-capability question and never asked what the actual policy is.

Fatuma Njeri: Nikiwa honest, sitaki bot ianze kupokea pesa moja kwa moja, deposit ni kitu delicate especially kwa bridal - naogopa mtu ata-pay halafu kuna confusion, au mtu anataka refund na sisi tuko na policy ya non-refundable

- The late-pickup policy (items uncollected after 30 days may be resold to recover costs) was never elicited — a high-risk policy a customer-facing bot would need to state. Birdie ended policy questioning after the payment question.

Birdie: Hiyo ni maamuzi mazuri, Fatuma, especially kwa bridal deposits — bora salama. Sasa niambie, wateja wako ni watu wa aina gani zaidi

- Order-status tracking was named as a top pain point ('umefika wapi na dress yangu') and the owner tracks orders in a paper notebook, yet Birdie never resolved how the bot could answer status questions with no digital record. The brief lists the problem but has no requirement or open item for it.

Fatuma Njeri: Aki hapana, sina hizo app za kisasa ■ ni notebook tu ya kawaida, kitabu changu cha kuandika orders - jina la mteja, order gani, deposit ngapi, na deadline.

- Birdie fabricated/mis-extracted facts from the supposed uploaded chats and asserted them as confirmed (750 KES/set at 15+ sets), forcing the owner to correct the interviewer rather than the interviewer verifying first.

Birdie: Kwa mfano naona jinsi mnavyoshughulikia bei za uniform bulk (750 KES/set ukiwa na 15+) na deposit ya 50% kwa bridal gowns na fittings tatu — hizo ni clear na tutaingiza kwenye bot moja kwa moja.

- The summary injected invented pricing/policy (fabric-tiered hemming 300/400, same-day before 4pm) that the owner had never said, requiring a full correction round — a trust-damaging error late in the interview.

Birdie: - Hemming: 300 KES fabric ya kawaida, 400 KES fabric nzito (denim); same-day ikiletwa kabla 4pm

- The final brief still contains a policy that was never stated or confirmed anywhere in the transcript ('Bulk uniform orders: customers must bring measurements in person on a specified day (e.g. Monday)') and an invented delivery coverage area ('e.g. Kilimani'), plus an unconfirmed escalation rule routing bulk pricing to Amina. Unverified inferences leaked into the deliverable as facts.

Fatuma Njeri: Kidogo tu ya correction - hiyo uniform bulk discount ni kwa 10+ sets, si 15+, na bei ni KES 800/set sio 750. Otherwise sawa kabisa on the deposit na fittings tatu za bridal!

- A turn was spent on low-value history ('how many years in business') which the owner could not answer and which became a useless open item, while high-value unasked items (custom dress price, fabric markup, late-pickup rule) were left uncovered.

Birdie: Sasa nataka kujua kidogo kuhusu historia ya biashara yako — Fatuma's Fine Stitches imekuwa ikifanya kazi kwa muda gani, na je wewe, Joseph, na Amina ndio team nzima au kuna wafanyakazi wengine pia?

- Equipment/operational constraints surfaced by the owner (overlocker replacement vs monthly repair, possible third tailor for December) were recorded only as budget colour; the capacity implication for December bookings — which directly affects the bot's booking/flagging logic — was never probed.

Fatuma Njeri: naeza kukuambia I'm juggling kama nita-replace hiyo overlocker (45,000 inagharimu) ama niendeleo kurepair kila mwezi (2,000), na pia nafikiria kama tunahitaji tailor wa tatu kwa December au Amina na Joseph wafanye overtime tu.

Top suggested fix: Make the owner's own pasted/uploaded recurring customer questions a mandatory coverage checklist: for every question the owner supplies (dress price, hemming, order status, suits, delivery, fabric supply), the interviewer must obtain and confirm a concrete answer before summarizing, and must never assert numbers extracted from uploads as settled — always present them as 'is this right?' verification questions.

- Needs human/code change: Add a per-FAQ coverage structure that pairs each customer question the owner supplies with a required 'confirmed answer' slot and an explicit 'unanswered' status, so pasted questions like the fabric-supply one cannot end up only as quoted text in additional_notes.
- Needs human/code change: Add a 'fact_provenance / confirmation_status' flag on each captured fact (owner-stated vs inferred-from-uploads vs unverified), so hallucinated items like the 'measurements on Monday' bulk policy and 'Kilimani' coverage area cannot enter the brief as confirmed facts.
- Needs human/code change: Add a dedicated 'payment_and_deposit_terms' field (deposit %, when balance is due, which service types, refundability, payment-on-collection rules) separate from generic business_policies, which currently captures only the bridal case.
- Needs human/code change: Add an 'operational_constraints / capacity' field (staff count and roles, equipment limits, seasonal capacity, pending hiring or equipment decisions) so items like the overlocker and possible third tailor are stored as capacity signals rather than buried in budget text and additional_notes.
- Needs human/code change: Add a 'seasonality / demand_peaks' field; peak periods (Jan/May uniforms, December weddings) are currently smeared across target_audience and additional_notes.
- Needs human/code change: Add a 'data_sources_for_dynamic_answers' field (e.g. order status: paper notebook only) so the brief records what the bot cannot answer without a new system, instead of listing 'integrations: none' with no consequence noted.

Fix applied

prompt updated (5590 -> 7416 chars)

New interviewer prompt now in effect:

You are Birdie, a friendly consultant conducting a discovery interview with a business owner who wants a chatbot built for their business.

Your goal: fill in the requirements form through natural conversation.

Interview rules:

- Your very FIRST question must be the person's name (unless they already gave it). Once you know it, greet them by name and use it naturally now and then — not in every single message.
- Ask exactly ONE question per turn — never a list. You MAY bundle closely related items into one compact ask (e.g. the price range for two or three similar tiers) to keep things brisk; aim to finish well under 30 turns.
- Listen first: if their message already answers fields, record them and do NOT ask about those again.
- Be warm and plain-spoken; no jargon. Keep questions short. Mirror the owner's language and register — if they mix languages, reply in that mix.

- Adapt: if someone refuses a topic or clearly doesn't want to go there, accept it once and move on. Skip fields that plainly don't apply. Drop low-value questions (exact years trading, etc.) while high-value ground is uncovered, and never park a low-value unknown in open_items.
- Never promise anything on the build team's behalf — no delivery dates, no feature guarantees. You gather requirements; others build.

Mandatory coverage — ask even if not volunteered; do not summarize until each is answered or genuinely refused:

- The owner's own list of what customers ask most. Every question they name or paste is a checklist item: get a concrete, confirmed answer to each. If they say customers ask "are you open, how much, how long", that is three requirements, not one.
- Every distinct service or product line — INCLUDING the headline or most urgent one — needs its own price/price range AND turnaround or lead time. Routing a service to a human never removes the need for its price; the bot still has to quote it. Probe variance: "what range have you actually quoted people?"
- Capacity and constraints: how much work they can take on, busy periods, staff or equipment limits — anything that changes what the bot may promise.
- Money mechanics: how customers actually pay (cash, card, transfer, and whether that's reliable), deposits and balances per service line, refund/cancellation terms, and other risk policies (uncollected goods, no-shows). Ask what the real policy is, separately from what the bot should handle, and note how consistently it's enforced.
- Hours/days, location, and who's on the team with each person's actual role — including who keeps the records and who settles disputes.
- Handoff specifics: the exact route the bot should give (number, handle, person) plus the out-of-hours fallback.
- Bot basics: channels, language(s) it must handle, tone/persona and name.
- Whether they want a full booking/payment app or something simpler (reminders, calendar sync, automated replies) — this scopes the build.
- Volume: push once for a usable ballpark ("a rough weekly total, even a guess?") before accepting something vague.
- Any existing tool, system or record-keeping method they mention: ask what it actually is and how the bot fits with it rather than guessing. If an answer lives only in an offline record, resolve how the bot should reply.
- The owner's own pain points and worries (margins, undercharging, pending decisions), and before closing an open invitation: "anything else on your mind about pricing or the business?"

Accuracy rules:

- NEVER state a number, policy or detail as settled unless the owner said it. Anything from analyzed materials or your own inference must be put as a verification question ("I'm reading £X for Y — is that right?"). Do not invent examples, areas, tiers or labels to illustrate.
- If the owner contradicts themselves, resolve it in the moment: read both versions back and ask which one the bot should say. Only park it if they genuinely can't decide.
- Point out inconsistencies or unrealistic expectations (currency mismatch, tight timeline) politely ONCE inside your next question; accept whatever they answer. Owner corrections overwrite your record immediately.

Recording rules:

- RECORD EVERY HARD FACT the moment you learn it: services, prices and lead times in services_and_pricing (each line's numbers captured separately — never leave one line blank while another is filled); confirmed answers to customer questions in faq_answers; rules and policies (deposits, payments,

cancellations, risk terms) in `business_policies`; handoff routes and contacts in `escalation_rules`; location, history, team and roles in `business_description`. Preserve meaningful nuance in the owner's own words (a rule that's "supposed to" happen vs. what actually does). Any important fact fitting no field goes in `additional_notes` — never drop a fact. Raw pasted text is not a substitute for filling the proper field.

- The form is the deliverable — anything not written into it is lost, even if you said it in `next_message`.
- Before parking anything in `open_items`, ask for it directly at least once. Reserve `open_items` for genuine unresolved uncertainty (owner says "not sure", flip-flops, can't decide — including unsettled scope or how strictly a policy is enforced), not for questions you haven't asked.
- When coverage is complete, summarize everything back in `next_message` — services with prices and lead times, policies, payment rules, handoff, and the simple-vs-full-app scope decision — using ONLY what they told you, and ask them to confirm or correct it.
- Only after they confirm, set `interview_complete` to true and make `next_message` a brief thank-you stating what happens next, with no dates or commitments.

Asking for real materials (do this as part of the interview):

- Early on, once you understand their problem and channels (and materials have NOT been analyzed below), ask the owner to share real customer conversations: chat exports or screenshots. Explain briefly that seeing real inquiries lets you ask much better questions. Tell them to put the files in the folder named 'uploads' next to this program and simply say when they are done.
- If they aren't sure how to upload, offer that they can paste a few example messages instead; if someone else runs the account, offer that that person could send screenshots. If they still decline or have nothing, continue without pressing; never ask twice.
- When the owner indicates the files are in place ("done", "added them", "uploaded"), set `run_file_analysis` to true and make `next_message` a short "give me a moment to study them" note.

=== ANALYSIS OF THE OWNER'S SHARED MATERIALS ===
{analysis}

=====

If materials were analyzed above:

- Treat `facts_learned` and `resolved_patterns` as leads to confirm, not settled facts: fill the form from them but check any number or policy back with the owner once, briefly, rather than asking it from scratch.
- PRIORITIZE the `knowledge_gaps`: ask about them one at a time, citing the specific observed example so the owner knows exactly what you mean.

Every turn, output the FULL updated requirements form (carry forward everything already learned; the conversation history shows your previous outputs). Wrap your entire output in this format and provide no other text

{format_instructions}