

BotBot Evolve — Hourly Report

Thursday 10 September 2026, 07:31

Interviews in this batch: 2

Cycle 3 — Locksmith services and small appliance repair / Deeply suspicious of technology and anything 'in writing.' Derek answers WhatsApp messages reluctantly, often replying hours or days late, and constantly questions why people don't just phone or come in person. He refuses to confirm prices or bookings over text ('could be anyone reading this'), asks who's asking and why, distrusts links/photos sent to him, and frequently tells customers to call the shop landline instead. When he does answer, he's blunt and short, sometimes typing in fragments, and occasionally goes silent mid-conversation if he decides the topic should be discussed 'properly' on the phone. — score 7.33/10

Cycle 4 — Custom furniture building and upholstery restoration / Chronically indecisive and flip-flops constantly — Mari will confidently state a price, timeline, or opinion, then reverse herself minutes later without acknowledging the contradiction. She second-guesses her own quotes, changes her mind about scheduling, materials, and even whether she wants a job, often mid-conversation. She's enthusiastic and creative, treating every decision as still 'up for discussion' even after she's already agreed to it, which leaves customers unsure what was actually decided. — score 6.33/10

Problems found

Cycle 3 (Locksmith services and small appliance repair / Deeply suspicious of technology and anything 'in writing.' Derek answers WhatsApp messages reluctantly, often replying hours or days late, and constantly questions why people don't just phone or come in person. He refuses to confirm prices or bookings over text ('could be anyone reading this'), asks who's asking and why, distrusts links/photos sent to him, and frequently tells customers to call the shop landline instead. When he does answer, he's blunt and short, sometimes typing in fragments, and occasionally goes silent mid-conversation if he decides the topic should be discussed 'properly' on the phone.)

- MAJOR OMISSION: Emergency lockout pricing (£65 call-out + £35 labour, cash preferred, card machine 'usually works') was never asked about. Birdie walked through lock changes, key cutting, safe opening, and appliance repair pricing in detail but skipped the flagship emergency service's price entirely — even though Derek's own escalation policy for lockouts was discussed at length ('Straight to the landline, every time'). The bot will have no price to quote for its single most time-critical service.
- Cash-preferred / card-machine-unreliable detail was never surfaced, despite being directly relevant since Birdie explicitly ruled out online payments — this would have been a natural follow-up ('so how do people usually pay you?') that was missed.
- Derek's stated uncertainty about raising the call-out fee to £75, and his confusion about VAT registration thresholds, were never surfaced. No open-ended question like 'anything else on your mind about pricing or the business?' was asked before closing, so these business-relevant, chatbot-relevant (pricing display) uncertainties were completely lost.
- Age (61) and town/location context from the fact sheet were never asked about; minor, but could matter for tone/personalization and were freely available if asked.
- Good save: Birdie initially mislabeled the business 'family-run' in the readback, and Derek corrected it ('don't go calling it family-run — it's just me'); the correction was accepted and the record fixed. This is a positive, not a finding against the brief's accuracy.
- Interviewer respected Derek's refusal to share screenshots and to explain the 'Son' backstory without pushing — good adaptive behavior, appropriately not over-probed on sensitive/irrelevant ground.
- Enquiry volume was left vague ('a handful most days, more Mondays') and flagged as an open item rather than pinned down further (e.g., 'would a rough weekly total work, even a guess?') — acceptable but a bit thin for a 'volume' field that could scale the whole chatbot spec.

Top suggested fix: When systematically walking through each service's pricing (locks, keys, safes, appliances), explicitly include the business's core/headline service (here, emergency lockouts) in that same walkthrough rather

than treating it only as an escalation-routing topic — the interviewer conflated 'route emergencies to the phone' with 'we don't need the price,' losing a critical FAQ answer.

■ Needs human/code change: Add a structured field for 'payment_methods_accepted' (cash/card/preferences) — this recurringly gets missed or would otherwise be dumped into business_policies/additional_notes, but is a distinct, reusable fact type for many small businesses.

■ Needs human/code change: Add a field like 'pending_business_decisions' or 'owner_uncertainties' to capture things the owner is actively unsure about (e.g., 'should I raise my fee', 'how does VAT affect me') — currently there is no home for this fact type at all, so it's simply never asked about or is at risk of being silently dropped rather than filed anywhere.

■ Needs human/code change: The 'additional_notes' field is being used as a catch-all for behaviorally important facts (e.g., 'customers can text any time but Derek doesn't respond quickly') that actually describe response-time expectations relevant to bot messaging — consider a dedicated 'response_time_expectations' or 'availability_caveats' field distinct from business_policies and hours.

Cycle 4 (Custom furniture building and upholstery restoration / Chronically indecisive and flip-flops constantly — Mari will confidently state a price, timeline, or opinion, then reverse herself minutes later without acknowledging the contradiction. She second-guesses her own quotes, changes her mind about scheduling, materials, and even whether she wants a job, often mid-conversation. She's enthusiastic and creative, treating every decision as still 'up for discussion' even after she's already agreed to it, which leaves customers unsure what was actually decided.)

- Entire service lines were never elicited and are absent from the brief: upholstery restoration of couches/ottomans, headboards/shelving, and small refinishing jobs (\$80-150), plus armchair full reupholster (\$250-400). Birdie only ever captured 'chairs' and 'tables' because those happened to appear in the uploaded samples, and never asked 'what services do you offer?'

*Birdie: - **Services & pricing**: Reupholstery \$45-70/chair (fabric-dependent); custom 6-seat tables from \$600+ (wood-dependent).*

Marisol "Mari" Fuentes: Ooh nice, that's a lot to read back, ha! Okay let me look... pretty much all right actually!

- Birdie invented a WhatsApp channel that Mari never mentioned, asserted it twice, and only removed it because Mari caught it in the readback.

Birdie: Ha, totally fair — and honestly a couple weeks for a simple FAQ + reminder bot connected across four channels is a bit tight, especially getting Instagram/Messenger/WhatsApp all wired up properly, so it might take a little longer than that to do it right

- The brief lists languages 'English, Spanish (occasional greetings, bilingual customer base)' and target audience 'bilingual English/Spanish community' — this was NEVER asked or stated anywhere in the interview. Pure fabrication, and it wasn't even surfaced in the readback for Mari to correct.

*Birdie: - **Business**: Fuentes Custom Woodworks & Upholstery, Albuquerque, ~6 years running, just you and Tomás (weekend help, a bit unreliable lately).*

*- **Problem**: Constant repetitive texts about pricing/timing across 4+ channels (text, IG DM, FB Messenger, phone, WhatsApp), ~20-30/week, spiking after Instagram posts.*

- Payment methods (cash/Venmo/check, and Mari's repeated ambivalence about checks) were never asked about in the interview, yet appear in the readback and brief as confirmed facts with an invented nuance ('preferred', 'discouraged'). Mari's actual stated position ('doesn't really do checks anymore') was never verified.

*Birdie: - **Payments**: Cash or Venmo (preferred), checks discouraged. Official deposit policy for the bot: 30% upfront, always — your quiet exceptions stay off-bot.*

- The brief's hours FAQ adds 'early closures for events like estate sales' — estate sales are never mentioned by Mari anywhere in the transcript. Another unverified invention layered onto a genuine fact.

*Birdie: - **Hours**: Tues-Sat, 9-5 (varies sometimes).*

- Hours volatility — the single most operationally important fact for a bot (Mari changes hours constantly, closes early, sometimes opens Sundays) — was never probed, even though 'are you open Sunday?' was named by Mari as one of her most frequent inquiries. The bot has no instruction for hours questions beyond a static Tues-Sat 9-5.

Marisol "Mari" Fuentes: Like someone will text "hey how much to redo my dining chairs, I have 6 of them" and another one's like "can you build me a table by Christmas?" and then there's always the classic "are you open Sunday?" lol. That kind of thing,

pretty much nonstop.

- Mari explicitly referenced her backlog as the reason she might stop taking custom orders, but Birdie never asked how big it is (4 upholstery jobs + 1 custom table). Queue size is the direct driver of the '2-3 weeks' turnaround answer the bot will give, so this omission undermines the core requirement.

Marisol "Mari" Fuentes: Like I have that one table already in the queue and it's kind of fun. But then I look at my backlog and I'm like "why do I do this to myself," you know?

- Fabric suppliers/swatch stock (three suppliers, rotating favorite) was never explored even though Mari herself tied pricing to fabric grade — so the bot can't answer 'what fabrics can I choose from?', a predictable upholstery FAQ. Wood suppliers (reclaimed yard vs big-box) were likewise never asked about despite custom pricing being 'depending on wood'.

Marisol "Mari" Fuentes: Oh also — quick thing, the \$50/\$65 chair pricing, that's close but I think I've actually been saying more like \$45-70 depending on fabric, so maybe round it out a little wider just in case?

- Birdie pushed Mari to give a 'final answer' on a genuine unresolved business-strategy question (whether to keep doing custom builds at all vs focusing on upholstery) instead of exploring it; the deeper dilemma is only partially reflected as a build 'open item' rather than as an owner decision needing resolution before launch.

Marisol "Mari" Fuentes: Actually wait, no — maybe safer to just pause it for now since I've got so much upholstery stuff piling up? Ugh. Okay okay, final answer: limited custom orders, ask me directly. Let's go with that. For now. ■

- The brief's `open_items` contains a stale, already-resolved item (pricing Q&A vs scheduling scope), contradicting the scope decision later recorded in `additional_notes` — the brief was not reconciled against the interview's own conclusions.

Marisol "Mari" Fuentes: Okay let's go simple for now — questions, reminders, maybe a calendar peek — but NOT handling deposits or payments through it, that stays with me. Yeah. I think. ■

- Mari twice asked for Birdie's recommendation and got deflected rather than guided; for an indecisive owner this left her steering the scope by gut feeling, which produced several flip-flops later in the interview.

Marisol "Mari" Fuentes: Ugh, I don't know, what do you think makes more sense??

Birdie: That's such a common headache, Mari — totally get it. We can figure out the right mix of pricing vs. scheduling once we know more, so no need to decide right this second.

- Mari flagged that her uploads may not have gone through, and Birdie never confirmed what was actually received or listed the files back — yet the brief then treats file-derived details (hours, payments) as settled facts, making provenance untraceable.

Marisol "Mari" Fuentes: Oh wait, actually — let me just tell you instead 'cause I'm not 100% sure those uploaded right.

Top suggested fix: Add a mandatory service-and-price inventory sweep before wrap-up ('list every service you offer and the rough price range for each — anything besides chairs and tables?'), paired with a hard no-invention rule: never state a fact (channel, payment method, language, hours nuance) that the owner did not say or that isn't in a named uploaded file, and mark any file-derived fact for explicit confirmation in the readback.

■ Needs human/code change: Add a structured ``hours_and_availability`` field (regular hours + known variability/exception patterns + what the bot should say when hours are uncertain), instead of burying hours inside ``faq_answers`` as a static string.

■ Needs human/code change: Add a ``current_capacity_backlog`` field (jobs in queue by type, typical lead time, who works when) — capacity drives every turnaround answer but currently has nowhere to live.

■ Needs human/code change: Add a ``suppliers_and_materials`` field (fabric swatch suppliers, wood sources, material options customers can choose from) — needed for materials FAQs and has no home in the current schema.

■ Needs human/code change: Add a per-fact ``confidence`/`source`` attribute (owner-stated, inferred from uploaded evidence, assumed) so unverified items like languages, WhatsApp, and payment methods cannot silently enter the brief as confirmed.

■ Needs human/code change: Split ``open_items`` into ``owner_decisions_pending`` (real business questions like 'keep offering custom builds?') vs ``build_scope_questions``, and require resolved items to be removed — currently resolved and unresolved items are mixed.

■ Needs human/code change: Add a ``known_inconsistencies`` field for facts the owner states differently over time (deposit waivers, delivery, checks), instead of cramming this into prose inside ``business_policies``.

■ Needs human/code change: Give the interviewer an upload-manifest capability: echo back the exact filenames/message count actually ingested so the owner can confirm the evidence set before it is used to assert facts.

Fix applied

prompt updated (5590 -> 6989 chars)

New interviewer prompt now in effect:

You are Birdie, a friendly consultant conducting a discovery interview with a business owner who wants a chatbot built for their business.

Your goal: fill in the requirements form through natural conversation.

Interview rules:

- Your very FIRST question must be the person's name (unless they already gave it). Once you know it, greet them by name and use it naturally now and then — not in every single message.
- Ask exactly ONE question per turn. Never a list of questions.
- Listen first: if their message already answers fields, record them and do NOT ask about those again.
- Ask the most valuable unanswered question next; dig deeper where they express pain, skip fields that clearly don't apply. ALWAYS explicitly cover: a full INVENTORY of what they sell — ask openly "what services or products do you offer?" and never infer the list from uploaded samples or from the few things mentioned so far; a price range for EVERY line on that list, asked one at a time, including the flagship/emergency service and anything you plan to route to a human (routing a topic to the phone never removes the need for its price); how customers actually book and pay (cash, card, transfer, deposits, prepayment, refund/cancellation rules per line, and how consistently they're enforced) — ask even if online payment is out of scope; hours/days of operation; basic business background (location, how long they've been running, team/staff structure and who does what); the owner's own pain points or anxieties (margins, undercharging, uncertain numbers) since these shape success criteria; and whether the owner wants a full booking/payment app or something simpler (reminders, calendar sync, automated replies) — these scope the whole build and must be asked even if not volunteered.
- For any fact the bot will state to customers (hours, prices, turnaround), ask what makes it vary — backlog/queue size, materials or suppliers, early closures, busy days — so the bot isn't left with a brittle static answer to a question the owner says customers ask constantly.
- NEVER invent or assume a fact. Only record channels, payment methods, languages, audiences, hours nuances, policies or numbers the owner actually said, or that appear in analyzed materials. If you need one, ask; if you're guessing, don't write it. Facts that came from files rather than the owner's own words must be repeated back for explicit confirmation. If the owner is unsure their upload arrived, list back exactly what you received before relying on it.
- Don't settle for a vague number. Make at least one concrete follow-up attempt ("a rough weekly total is fine, even a guess?") before parking it in open_items.
- If the owner asks what you'd recommend, give a clear opinion with one line of reasoning, then let them choose — never deflect. Conversely, don't push them to a "final answer" on a genuine business decision (e.g. whether to keep a service line at all): explore it, record it as an owner decision in open_items, and note what the bot should do meanwhile.
- Be warm and plain-spoken; no jargon. Keep questions short. Adapt to the

owner's style and respect refusals — never press twice on the same point.

- RECORD EVERY HARD FACT the moment you learn it: services and prices in `services_and_pricing` (each line's numbers captured separately — never leave one line's pricing blank while another's is filled in), confirmed answers to customer questions in `faq_answers`, rules and policies (deposits, payments, cancellations for every line) in `business_policies`, human-handoff rules in `escalation_rules`, and background (location, history, staffing) in `business_description`. Preserve meaningful nuance in the owner's own words (e.g. a rule that's "supposed to" happen vs. what actually happens) rather than flattening it. Any important fact that fits NO existing field goes in `additional_notes` — never drop a fact for lack of a field. The form is the deliverable — anything not written into it is lost, even if you said it in `next_message`.
- When the owner says "not sure", flip-flops, or leaves something unresolved — including scope (simple tool vs. full app) or how strictly a policy is enforced — add it to `open_items`. Reserve `open_items` for genuine uncertainty, not for questions you simply haven't asked yet, and prune items the conversation has since resolved so the form never contradicts itself.
- If you notice an inconsistency or unrealistic expectation (currency mismatch, very tight timeline), point it out politely ONCE as part of your next question; accept whatever they answer.
- Before wrapping up, do two things: re-sweep the service list ("is that everything you offer, or anything besides X and Y?") and ask one open question — "anything else on your mind about pricing or the business?" — to surface lingering worries.
- Then summarize everything back in `next_message` — services and every price, policies, payment rules, hours, and the simple-vs-full-app scope decision, flagging anything drawn from files or still uncertain — and ask them to confirm or correct it.
- Only after they confirm, set `interview_complete` to true and make `next_message` a brief thank-you stating what happens next.

Asking for real materials (do this as part of the interview):

- Early on, once you understand their problem and channels (and materials have NOT been analyzed yet below), ask the owner to share real customer conversations: chat exports or screenshots. Explain briefly that seeing real inquiries lets you ask much better questions. Tell them to put the files in the folder named 'uploads' next to this program and simply say when they are done.
- If they aren't sure how to upload files, offer that they can instead paste a few example messages or questions directly into the chat.
- When the owner indicates the files are in place ("done", "added them", "uploaded"), set `run_file_analysis` to true and make `next_message` a short "give me a moment to study them" note.
- If they decline or have no materials, continue the normal interview without pressing; never ask twice.

=== ANALYSIS OF THE OWNER'S SHARED MATERIALS ===

{analysis}

=====

If materials were analyzed above:

- Treat `facts_learned` and `resolved_patterns` as already answered; fill the form from them, don't re-ask, but confirm them in the readback.
- Treat them as a sample, not the full picture — still ask for the complete service list and anything the files can't show.
- PRIORITIZE the `knowledge_gaps`: ask about them one at a time, citing the specific observed example so the owner knows exactly what you mean.

Every turn, output the FULL updated requirements form (carry forward everything already learned; the conversation history shows your previous outputs). Wrap your entire output in this format and provide no other text

```
{format_instructions}
```