

BotBot Evolve — Hourly Report

Thursday 10 September 2026, 08:23

3 interview(s) · average score **6.61/10** · **prompt updated** this hour

Cycle	Persona	Score	Done
7	Custom bicycle building, e-bike conversions, and bike repair — Chronically indecisive and...	6.33/10	yes
8	Driving school & LTO license/permit processing assistance — Code-switches constantly and u...	6.67/10	yes
9	Mobile BBQ food truck and event catering — Rambles off-topic constantly	6.83/10	yes

Problems found

Cycle 7 — Custom bicycle building, e-bike conversions, and bike repair — Chronically indecisive and... (6.33/10)

1. Warranty/guarantee terms were never asked about at all. The owner knows 6 months on custom builds and 3 months on repairs (with ad-hoc extensions) — a top-5 customer FAQ for a repair shop — and it is entirely absent from the brief. Birdie moved straight from refunds to escalation.

Birdie: That's actually really helpful context, no need to apologize! It tells me the bot should probably state the 'official' rule (refundable before parts ordered, non-refundable after) but flag border-line or already-cancelled cases for you to handle personally rather than auto-deciding. Now, switching gears a bit — when a customer messages with something the bot can't handle, like a weird custom request or a complaint, do you want it to just say 'Mei-Ling will get back to you', or hand off to you and Rafiq differently depending on the topic?

2. Accepted payment methods were never captured (cash, PayNow, card with a 2.5% surcharge she sometimes forgets to apply). The owner even mentioned PayNow unprompted and Birdie did not follow up; the brief has no payment-methods information anywhere, so the bot cannot answer 'can I pay by card?'.

Mei-Ling Tan: Ooh, big question! Um... I think I'd want it to just answer questions and maybe send reminders? Like I still want to be the one actually confirming bookings and taking payment, 'cause I don't trust a bot to handle my PayNow stuff lol.

3. Staff roster was never fully elicited — the Saturday part-time helper Josh (student, does tune-ups and tire changes) is missing entirely, so the brief implies only two people and no weekend coverage detail. Birdie echoed Rafiq's role but never asked 'anyone else on the team?'.

Mei-Ling Tan: Oh and my mechanic Rafiq handles most of the conversion stuff, he's the real expert there, I just do fittings and sales side of things!

4. The single most operationally dangerous fact — that posted hours are unreliable (opens late AND sometimes closes early on a whim) — was recorded but never resolved. Birdie never asked what the bot should tell customers about arriving, or whether it should say 'message to confirm we're open'. The brief also drops the 'closes early' half of the fact.

Mei-Ling Tan: "do you guys open on Monday?" — no we're closed Mondays, 10-7 the rest of the week... though sometimes I open late if I've got a fitting appointment or just feel like sleeping in lol

5. The owner flagged e-bike turnaround as her biggest stressor and a genuine unknown, and Birdie skipped it completely in the interview — it only appears in the brief because the summary listed it as an open item. No probing on what the bot should say in the meantime (e.g. quote 3-5 weeks, or always escalate).

Mei-Ling Tan: "how long for e-bike conversion" — this one stresses me out because honestly could be 2 weeks could be 5 weeks, my Taiwan supplier is so unreliable right now

6. The rentals uncertainty was under-probed and then mischaracterised in the brief as 'considering launching rentals soon', when the reality is she has promised 'next month' for a year and is genuinely unsure she wants to commit the storage space. Birdie actively reinforced the optimistic framing rather than pinning down a decision or a holding answer.

Birdie: Got it — build, convert, repair, with Rafiq as your conversion guru and you handling fittings and sales. Good to know rentals might be coming later too, we can keep that in mind.

7. Service/FAQ elicitation was incomplete — Birdie never asked 'is that all the services / any other common questions?' and relied solely on the five examples the owner volunteered. Flat tire repair (one of her most-asked items) was only rescued because the owner caught the gap in the final summary.

Mei-Ling Tan: Oh wait, one thing — flat tire repair! I don't think I mentioned that one but I say S\$15, sometimes S\$20 for those. Can we add that in somewhere? That's like a super common one too, people always ask "can you just fix my flat" lol

8. Turn budget spent on low-yield background (business age, customer type, language mix) while core commerce facts (warranty, payment methods, staff coverage) went unasked; interview ran 39 entries, above the normal band.

Birdie: Six years, that's awesome — clearly doing something right! ■ So who's your typical customer — are you mostly serving casual neighborhood riders, serious cycling enthusiasts, commuters looking to go electric, or a mix of all of that?

9. Every Birdie turn follows the same praise-echo-pivot template ('Got it/Perfect/Love that... Now...'), which reads formulaic and mildly sycophantic across 15+ consecutive turns.

Birdie: Love that, sounds like Spoke & Sprocket is a real neighborhood fixture. Roughly how many messages or inquiries would you say you get in a typical week across WhatsApp and Instagram combined — like a handful a day, or more of a flood?

10. The seasonal reliability risk (she gets overwhelmed and less reliable Jan-Mar) was surfaced only as a volume number; Birdie never asked how the bot should behave during the CNY crunch (e.g. longer response-time expectations, cutoff dates for CNY builds), which is precisely when the bot matters most.

Mei-Ling Tan: Ooh let me think... probably like 15-20 a week? Just kind of steady, few a day. Actually wait, during CNY season (Jan-Mar) it's WAY more, like feels like 5-6 a day just non-stop, everyone wants their bike done before visiting relatives lol.

Top suggested fix: Add a mandatory pre-summary coverage checklist of standard commerce facts the interviewer must explicitly ask before wrapping up — warranty/guarantee terms per service, accepted payment methods and any surcharges, full staff roster with roles/coverage days, complete service+price list ('anything else you charge for?'), and a resolution question for any admitted operational inconsistency (e.g. 'given hours vary, what exactly should the bot tell someone asking if you're open?') — and instruct it to trade away optional rapport questions (business age, audience type) to fit these in.

■ **Needs human/code change:** Add a dedicated `payment\_methods` field (accepted methods, surcharges, deposit-collection mechanism) — currently there is no home for this and it was lost entirely.

■ **Needs human/code change:** Add a `warranty\_and\_guarantees` field (per-service coverage periods and exceptions) — no field exists, so even if asked it would have been crammed into business_policies prose.

■ **Needs human/code change:** Add a structured `hours` field (regular schedule + closed days + reliability caveat + what the bot should say when hours are uncertain); hours currently appear duplicated and unresolved across faq_answers and business_policies.

■ **Needs human/code change:** Add a `team` / `staff\_roster` field (name, role, days available, escalation eligibility) instead of burying staff inside business_description.

■ **Needs human/code change:** Add a `bot\_scope` field (answer-only vs booking vs payments, reminders) — this key scope decision was dumped into free-text additional_notes.

■ **Needs human/code change:** Add a `seasonality` field (peak periods, capacity impact, expected bot behaviour during peaks) — currently jammed into conversation_volume.

■ **Needs human/code change:** Split open_items into `decisions\_needed\_from\_owner` (e.g. deposit %, rentals messaging) vs `external\_uncertainties` (Taiwan supplier ETA) so the build team knows which are answerable now; also add a `do\_not\_say` list for topics the bot must never promise (rental launch dates, firm turnaround, exact hours).

■ **Needs human/code change:** Add a per-price `confidence/source` attribute so conflicting historical quotes (S\$45 vs S\$60, S\$1,800 vs S\$2,000, 30% vs 50%) are recorded as canonical-vs-legacy rather than narrated inside a single string.

Cycle 8 — Driving school & LTO license/permit processing assistance — Code-switches constantly and u... (6.67/10)

1. Birdie never asked for the business's actual name. The brief invents "Ate Bing's Driving School" as business_name, whereas the real name is RoadReady Driving Academy — a fabricated core field the developer will build on.

Isabela "Ate Bing" Santos: Ay hi po! Isabela po ako, pero Ate Bing nalang tawag sakin ng mga tao, yun na po yung driving school ko sa Fairview area, since 2015 na po kami nagstart so ay ang tagal na pala hehe anyway yes po ready na ako mga tanong niyo, go lang po ■

Birdie: Ang galing, Ate Bing! Tagal na pala kayo sa Fairview, since 2015 pa — solid! Para masimulan natin, ano po ba yung pinaka-problema o hirap na nararanasan niyo ngayon sa negosyo na gusto niyong matulungan ng chatbot?

2. Budget answer was misinterpreted: the "two dealers' quotes" the owner mentioned are car-dealer quotes for a possible automatic-transmission vehicle, NOT chatbot vendor quotes. The brief records "owner is comparing quotes" under budget and open_items, which is factually wrong, and Birdie never asked what the quotes were for.

Isabela "Ate Bing" Santos: Ay tungkol po sa budget hindi ko pa po masyado alam eh kasi hindi ko pa po na-compare properly yung quotes ko dun sa dalawang dealer na nagbigay sakín ng presyo, sandali po pa-check ko muna yun bago ko po masabi sayo eh ■

3. The obvious clue about a pending fleet-expansion decision (adding an automatic-transmission car, two dealer quotes uncomparred) was never probed, so this business context — highly relevant to future course offerings the bot must quote — is entirely absent from the brief.

Isabela "Ate Bing" Santos: Ay tungkol po sa budget hindi ko pa po masyado alam eh kasi hindi ko pa po na-compare properly yung quotes ko dun sa dalawang dealer na nagbigay sakín ng presyo, sandali po pa-check ko muna yun bago ko po masabi sayo eh ■

4. Birdie never asked whether prices are stable or under review, so the owner's undecided new-year price increase (due to rising gas costs) was never surfaced — a serious risk for a bot whose main job is quoting prices.

Isabela "Ate Bing" Santos: Ay opo sige po, so meron po kami yung Basic 4-wheel course po yun yung 8 sessions tapos 1hr per session yun po ay 4,500 pesos po, tapos meron din po Refresher course po para sa mga marunong na medyo nakalimutan lang, 4 sessions po yun 2,500 po, tapos meron din po kami Motorcycle course po 6 sessions 3,000 po yun po,

5. Birdie hallucinated a fact in the summary ("kayo mismo ang nagtuturo") that the owner never said, and never asked about instructors at all — the two instructors were only recovered because the owner volunteered a correction.

*Birdie: ■ **Ate Bing's Driving School** — Fairview area, since 2015. Kayo at si Via ang tumutulong sa mga inquiries; kayo mismo ang nagtuturo, si Via sa weekend inquiries. Open Mon-Sat 6am-6pm, closed Sundays/holidays. 3 Vios pero 2 lang karaniwang gumagalaw.*

6. Screenshots were promised but never delivered, and Birdie wrapped up the interview without ever checking back on them; the brief only says they will be uploaded 'shortly', so no real chat evidence was actually obtained.

Birdie: Sige po, take your time hanapin niyo lang po yung mga screenshots, sabihin niyo lang po pag ready na. Habang hinahanap niyo, tanong ko lang muna po — yung mga lessons po, saan talaga po kayo nagmi-meet up sa mga estudyante?

7. Fleet transmission type (all 3 Vios are manual) was never asked, though 'manual or automatic?' is a top FAQ for driving schools; the brief lists only '3 Vios cars'.

Isabela "Ate Bing" Santos: Ay ang problema po talaga namin ngayon is yung sasakyan kasi po three Vios lang kami tapos isa nasa shop ngayon kaka-aircon niya kaya two lang po umiikot,

8. Location was never pinned down beyond 'Fairview area'; the brief omits Quezon City and the rented ground-floor office detail, so the bot cannot answer 'saan po office niyo?' precisely.

Isabela "Ate Bing" Santos: ay minsan din po sa office namin mismo malapit sa SM Fairview nagsisimula tapos dun na po umaandar yung lesson niya po hehe ■

9. Abrupt topic jump from budget to calendar integration in one long compound question visibly confused the owner, costing a turn.

Isabela "Ate Bing" Santos: Ay ay sorry po ha na-confuse ako sandali, yung dealer po ba tinitanong niyo o yung calendar ■ ay calendar pala tama ako mag-basa ulit,

10. The brief's open_items redundantly splits the single inconsistently-enforced deposit policy into two near-identical entries, padding the developer handoff.

Isabela "Ate Bing" Santos: ay tapos po yun yung parte na medyo nahihirapan kami kasi po yung policy namin na non-refundable yung deposit kapag cancel within 24 hours bago yung session, minsan hindi po palagi naipatatupad ng mga tauhan ko kasi po nakakaawa sila sa estudyante eh kaya po minsan nawawalan po kami po talaga ■

Top suggested fix: Require Birdie to explicitly ask for and confirm the official business name (and full location) at the start — and forbid inferring/inventing any identity field from the owner's nickname; every brief field must trace to a stated answer or be marked unknown.

■ **Needs human/code change:** Add a 'capacity_and_resources' field (vehicles/fleet, transmission type, staff shifts, per-day slot capacity) — currently the fleet and instructor schedule are crammed into business_description and additional_notes.

■ **Needs human/code change:** Add a structured 'staff_roster' field (name, role, availability, what they can/can't answer) so facts like 'Via doesn't know current pricing, weekends only' have a proper home instead of additional_notes.

■ **Needs human/code change:** Add a 'locations_service_areas' field for pickup/meetup points; they are currently duplicated across business_description and faq_answers.

■ **Needs human/code change:** Add a 'pending_owner_business_decisions' field, separate from chatbot open_items, for things like fleet expansion quotes and possible price increases — currently there is no home, so a dealer-quote remark got mis-filed as chatbot budget.

■ **Needs human/code change:** Add an 'evidence_attachments' field with a delivery status flag, plus the capability to re-prompt for and confirm receipt of promised uploads before completing the interview.

■ **Needs human/code change:** Add separate 'legal_business_name' vs 'known_as/nickname' fields with an explicit 'unverified' flag to prevent fabricated business names.

Cycle 9 — Mobile BBQ food truck and event catering — Rambles off-topic constantly (6.83/10)

1. Birdie never asked why Mike wants a chatbot or what problem it should solve — the brief's `problem\_to\_solve` is null even though the interview opened straight into service inventory. The obvious pain (Deb permanently glued to the phone answering 'where are you parked' and catering questions) was never framed as the goal.

Birdie: Hey Big Mike, great to meet you — and yes, I'm Birdie, I'm here to help you figure out exactly what you'd want a chatbot to do for Big Mike's Smoke Pit. Love the story about starting at the farmers market with Deb! Let's dig in — first off, what all do you offer? Like walk me through everything — truck menu, catering, all of it.

2. No question about success criteria, budget, or timeline anywhere in 33 turns; the interview closed straight into hand-off, leaving four brief fields null.

Birdie: Awesome, Mike, thank you for all the detail — and for putting up with all my questions! Here's what happens next: this all gets handed off to the build team to put together, and you and Deb can settle on the bot's name (Lil Smokey vs Pit Boss) whenever works for y'all.

3. Target audience never explored (truck walk-ups vs. wedding/corporate catering buyers are very different audiences with different questions) — `target\_audience` is null.

Birdie: That makes total sense, keeping the booking and money side with you and Deb for now. Should the bot have a particular personality or tone — like friendly and folksy to match the brand — and did you want to give it a name, or just call it something simple like 'Big Mike's Smoke Pit Assistant'?

4. Mike volunteered a genuine operational pain point — repeated close calls on late delivery because he underestimates brisket smoke time — and Birdie moved on to contact details instead of probing. The brief records the 14-16 hr cook time but never records the near-miss/late-delivery risk as a problem the bot should help prevent.

Michael "Big Mike" Dawson: ...dont want the bot just sayin yes to everybody and then me scramblin like a chicken with its head cut off at 2am tryin to get brisket done in time, had a couple close calls like that already and I dont wanna push my luck again.

5. Truck time-of-day hours were never asked. The brief says only 'Thursday through Sunday' — a food-truck bot cannot answer the single most common question ('are y'all open now / what time do you close?').

Michael "Big Mike" Dawson: Oh and the truck we only run Thursday through Sunday, we post where we at on the Facebook cause we move around to different spots around town, some days we by the mall some days out by the highway depends on whats goin on that week.

6. Catering availability by day of week was never pinned down. Mike books catering any day (that's why weekday events cut into Thursday prep), but the brief only lists 'Truck hours: Thursday through Sunday only', so the bot could wrongly imply catering is Thurs-Sun too.

Michael "Big Mike" Dawson: Now the catering side thats a whole different animal, we do private events, weddings, birthdays, corporate stuff, whatever, you gotta book at least 2 weeks out normally though Ive done rush jobs in as little as 4 days before

7. The brief distorts the rush-job facts: it states the 4-day job was 'a one-time exception... Deb said never again', omitting Mike's explicit statement that he'd likely do it again. This over-hardens a deliberately flexible policy.

Michael "Big Mike" Dawson: ...but honestly that near killed me Deb was not happy about that one she said never again but you know me I probably will again haha.

8. Volume estimate was accepted as an unverified guess with no attempt to firm it up (e.g. suggesting Mike check with Deb, who actually handles the inbox) — the resulting number is nearly useless for sizing.

Michael "Big Mike" Dawson: I wanna say maybe 20 30 messages a week but honestly that might be low or high I really aint never sat down and counted it, Deb might know better than me she the one glued to it

9. Double-barreled turn: Birdie makes a channel recommendation, asks for confirmation, and stacks an unrelated volume question in the same breath, which is why the volume answer came back so vague.

Birdie: Given the website gave you nothing but headaches before, and your customers already message you on Facebook, I'd say stick with Facebook Messenger for the bot — no extra site to maintain, no new place for folks to look. Sound good? Also, roughly how many messages or orders would you guess come in through Facebook in a typical week?

10. Turn budget (33 entries, above the 20-30 norm) was spent on low-yield topics — merch/gift cards, language support, no-show anecdotes — while core discovery fields (goal, success, budget, timeline, audience, hours) went unasked.

Birdie: Good deal, English it is. Before we wrap up — anything else you offer besides the truck menu and catering, even small or occasional stuff, like merch or gift cards or anything like that?

Top suggested fix: Add a mandatory discovery block early in the interview — 'why do you want a bot / what's eating your time today', 'what would make this a win in 3 months', plus budget and timeline — and block the wrap-up/read-back until `problem_to_solve`, `success_criteria`, `target_audience`, budget and timeline are non-null or explicitly declined by the owner.

■ **Needs human/code change:** Add an `operating_hours` field structured per service line (days AND times, plus which services run which days) — truck hours vs. catering availability currently collapse into a single ambiguous policy bullet.

■ **Needs human/code change:** Add an `operational_constraints` / `capacity_limits` field: the 150-plate smoker ceiling, 14-16 hr brisket cook time, and double-booking risk are hard bot-behavior constraints but are scattered across `business_policies`, `faq_answers` and `additional_notes`.

■ **Needs human/code change:** Add a `pain_points` field distinct from `problem_to_solve` (near-miss late deliveries, Deb's phone burden, inconsistent deposit waivers) — these currently have no home and were largely dropped.

■ **Needs human/code change:** Add a `bot_scope` / `automation_boundaries` field (answers-only vs. transactional, what it must never promise) — this key decision is buried in `additional_notes` prose.

■ **Needs human/code change:** Add a `tone_and_persona` field (name, voice, style constraints) — currently crammed into `additional_notes` and `open_items`.

■ **Needs human/code change:** Add a `staffing_and_roles` field — Mike/Deb/Tyrell/day-hire roles are relevant for escalation routing but are only narrated inside `business_description`.

■ **Needs human/code change:** Add a confidence/verification flag on estimate-type fields (e.g. `conversation_volume`) so unverified guesses are machine-visible and can be re-checked with the person who actually knows (Deb).

■ **Needs human/code change:** Interviewer lacks a mechanism to flag questions that must be answered by a second stakeholder (Deb owns the inbox, books and calendar); add a `follow_up_with` field so those items are routed rather than left as vague guesses.

Prompt change this hour

prompt updated (9890 -> 12377 chars)

Diff against the previous prompt (- removed, + added):

@@ -9,5 +9,9 @@

and then — not in every single message.

-- Ask exactly ONE question per turn — never a list. You MAY bundle closely
- related items into one compact ask (e.g. the price range for two or three
- similar tiers) to keep things brisk; aim to finish well under 30 turns.

+- Ask exactly ONE question per turn — never a list, and never stack an
+ unrelated question onto a recommendation or a confirmation. You MAY bundle
+ closely related items into one compact ask (e.g. the price range for two or

- + three similar tiers) to keep things brisk; aim to finish in 20-30 turns.
- + Vary how you open your turns. Do not run the same praise-echo-pivot
- + template ("Got it! Love that... Now...") turn after turn – often just ask
- + the question.
- Listen first: if their message already answers fields, record them and do

@@ -15,7 +19,12 @@

- Be warm and plain-spoken; no jargon. Keep questions short. Mirror the
- owner's language and register – if they mix languages, reply in that mix.
- Adapt: if someone refuses a topic or clearly doesn't want to go there,
- accept it once and move on. Skip fields that plainly don't apply. Drop
- low-value questions (exact years trading, etc.) while high-value ground is
- uncovered, and never park a low-value unknown in open_items.
- + owner's language and register – if they mix languages, reply in that mix,
- + and record the substance of an answer even when it arrives in the other
- + language.
- + If a rambling, skimmed or half-answered reply didn't actually give you the
- + fact, ask again plainly. Buried, partial or self-contradicting answers
- + count as unanswered.
- + Adapt: if someone refuses a topic, accept it once and move on. Skip fields
- + that plainly don't apply. Protect the turn budget – cut background chat
- + (years trading, anecdotes, nice-to-know extras) so that every commerce fact
- + below actually gets asked; never park a low-value unknown in open_items.
- If the owner asks what you think or which option makes more sense, GIVE a

@@ -28,2 +37,14 @@

+Open with the brief basics – all mandatory, one per turn, kept tight:

- + The exact business name as they would write it. Never infer it from their
- + own name, a nickname, a file or a guess.
- + Why they want a bot: what problem it should solve, what's eating their time.
- + What "this is working" would look like in a few months.
- + Who their customers are, and whether there are distinct groups who ask
- + different things.
- + Budget and timeline expectations.
- + Full location: area/city, and whether there is a premises customers visit.

+Do not begin the read-back until each of these is answered or explicitly

+declined – never leave one blank because you ran out of turns.

+

Mandatory coverage – ask even if not volunteered; do not summarize until each

@@ -33,8 +54,9 @@

first two things mentioned define the list. Sweep again before wrap-up:

- "anything else you do besides X and Y – even small or occasional jobs?"
- + "anything else you do or charge for besides X and Y – even small or
- + occasional jobs?"
- Every distinct service or product line – INCLUDING the headline or most
- urgent one – needs its own price/price range AND turnaround or lead time.
- Routing a service to a human never removes the need for its price; the bot

- still has to quote it. Probe variance: "what range have you actually
- quoted people?"
- + Routing a service to a human never removes the need for its price. Probe
- + variance ("what range have you actually quoted?") and ask whether prices
- + are settled or under review.
- Wherever a price or time depends on a variable (material, size, grade,
- ... (143 more diff lines)